




Quickly debug VM failures in OpenStack

Chenli Hu
chhu@redhat.com



What happened ?

How do you know
that ?

How to debug the failures quickly ?

Solutions:

- ★ Search engine
- ★ Research manual/ doc/ code
- ★ Attempt more times
- ★

Experience

Auxiliary tools

Agenda

- Premise
- Know more about VM start implementation
- Know more about debug logs
- Debug journey
- Call for actions

Premise


- Based on RHOS (Red Hat OpenStack)
- Compute node with libvirt driver & Qemu Hypervisor
- Focus on VM operations

OpenStack Compute (Nova) drivers:


- libvirt
 - Qemu/KVM on x86
(fully supported)
- Hyper-V, VMware, XenServer,
xen via libvirt
- baremetal, docker, LXC via libvirt

<https://wiki.openstack.org/wiki/HypervisorSupportMatrix>

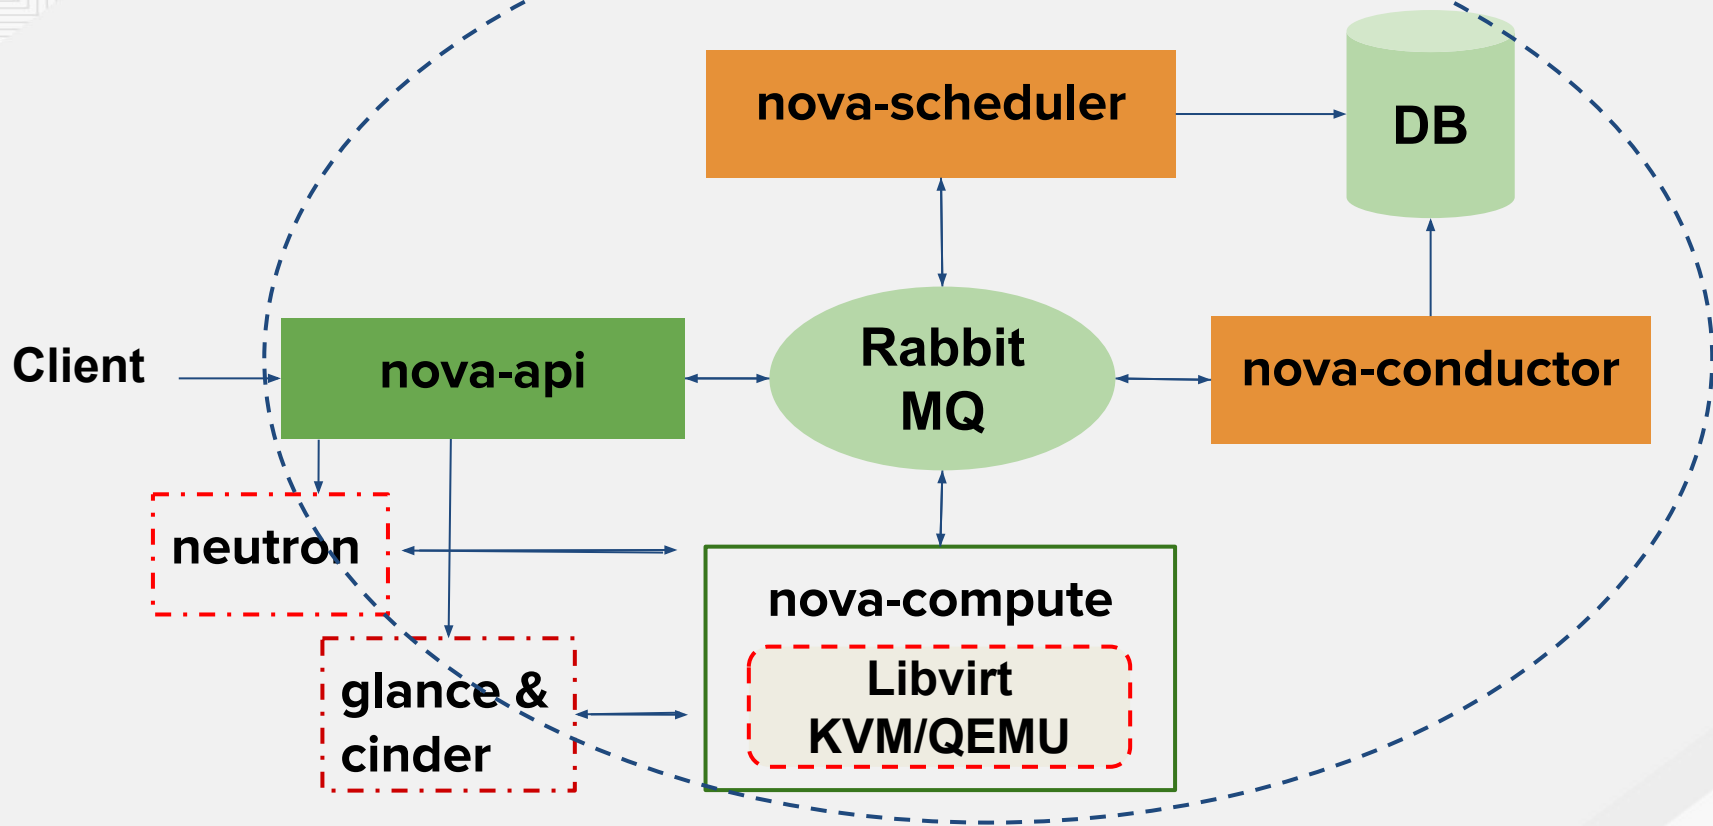
VM start failure



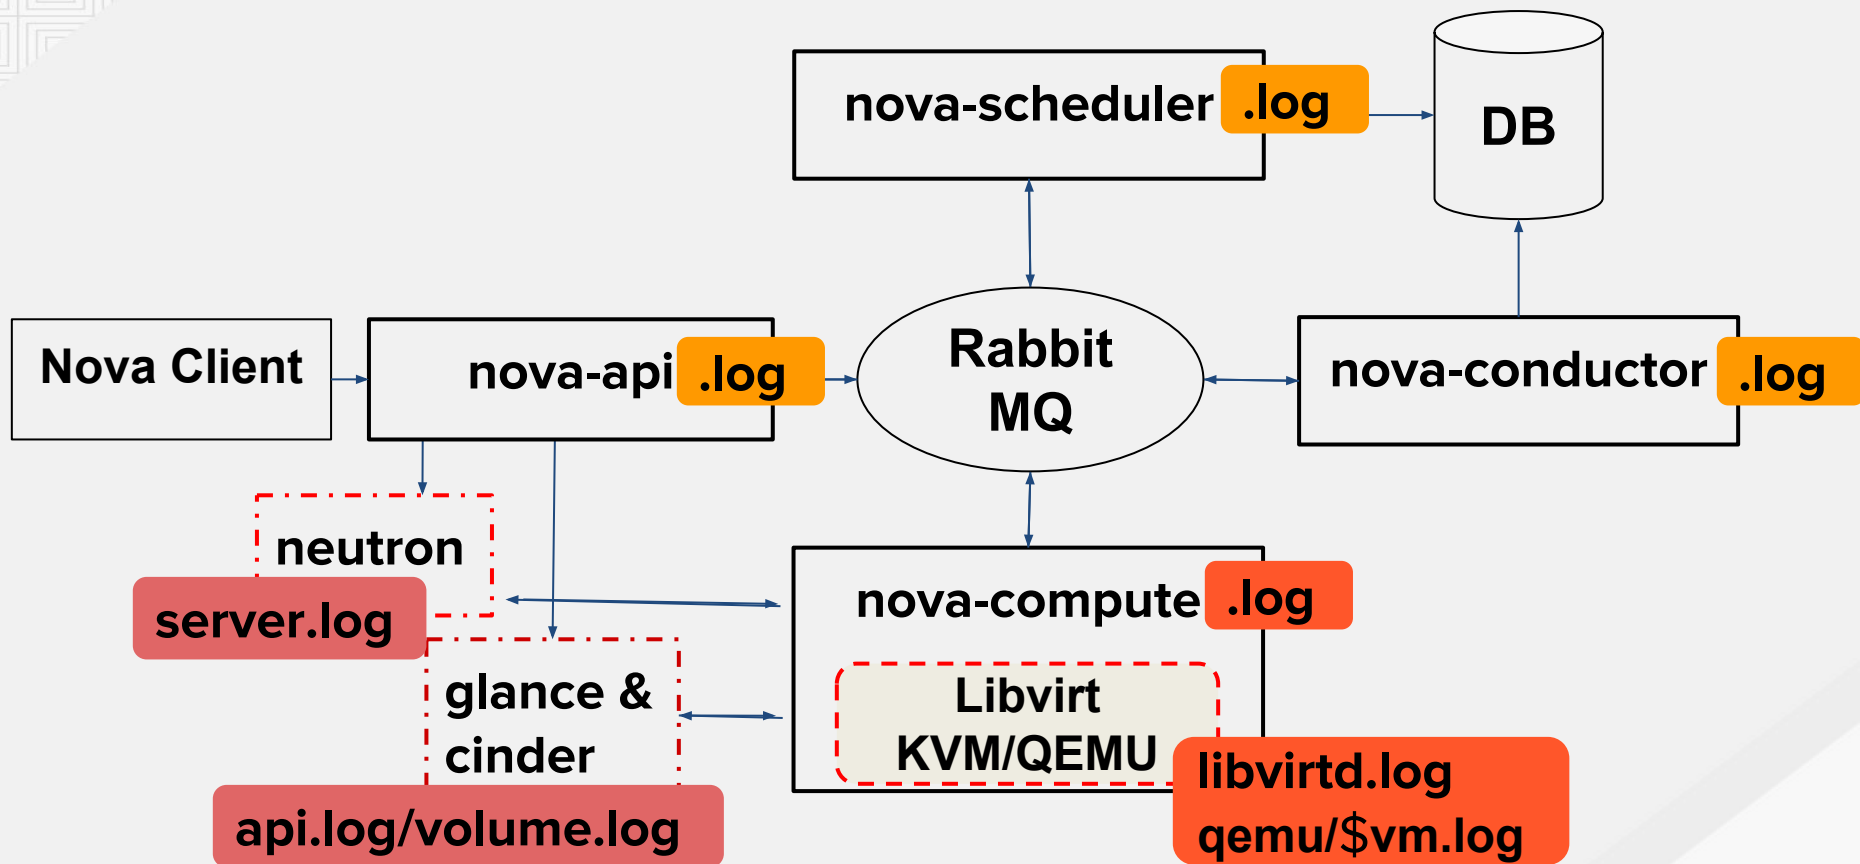
Where is VM running in OpenStack ?



How does VM start ?



VM start - Debug logs



VM start - nova-compute.log

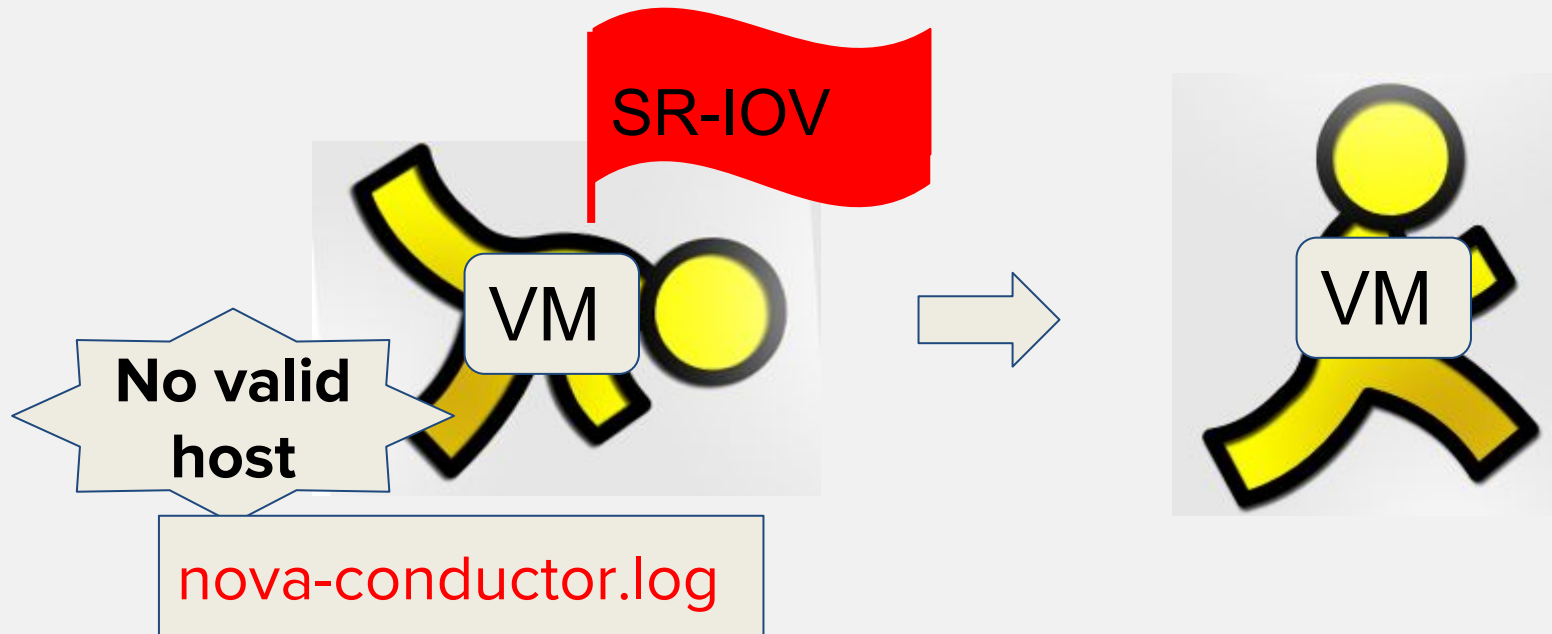
- start VM successfully:

```
nova.compute.claims ...Attempting claim: memory 1024 MB,  
disk 12 GB, vcpus 2 CPU.....  
nova.compute.claims...Claim successful  
nova.virt.libvirt.driver...Creating image  
os_vif ...Successfully plugged vif VIFBridge  
(...network=Network(...),plugin='ovs',port_profile=VIFPortProfileBase,  
preserve_on_delete=False,vif_name='tap00d8fa62-3f')  
nova.compute.manager...VM Started (Lifecycle Event)  
nova.virt.libvirt.driver [-] [instance:...] Instance spawned successfully
```


VM failure types:

- Lack of resource 
- Configuration problem 
- Not support 
- Product issue 
 - Qemu/ libvirtd service crash => ABRT

1 - VM start failure - Lack of resource



2 - VM live migration - Configuration problem




SELinux setting ? `#setsebool -P virt_use_nfs 1`

VM dir permission ?

NFS problem ?

3 - VM block migration - Not support




`nova-compute.log`

`virsh command`

Live Migration failure: internal error: unable to execute QEMU command 'migrate': this feature or command is not currently supported

4 - VM block migration - libvirtd/qemu crash



libvirtd.log

virsh command

ABRT

debug : qemuProcessHandleMonitorEOF:.....assuming
the domain crashed

Debug journey

nova
neutron, glance, cinder

Libvirt/ Qemu-kvm

Bug ?

Debug journey - settings in OpenStack

➤ VM related settings:

/etc/nova/nova.conf; image-update/create; flavor settings

images_type
(raw, qcow2, lvm, rbd,
default)

cpu_mode
vcpu_pin_set
reserved_huge_pages

max_concurrent_live_mig
rations
live_migration_uri
live_migration_tunelled
live_migration_permit_po
st_copy
.....

glance image-update
flavor settings:
hw_disk_bus
hw_qemu_guest
hw_scsi_model
hw_vif_model
.....

glance image-update
flavor settings:
hw_numa_nodes
hw_numa_memory
hw_numa_cpus.0/1
hw_numa_mem.0/1
.....

Debug journey - logs in OpenStack

➤ Debug settings:

log_dir=/var/log/nova

➤ Debug logs:

- VM life-cycle, migration
 - nova-api.log, nova-compute.log,
 - nova-conductor.log, nova-scheduler.log
- Images
 - glance/api.log
- Volume
 - cinder/api.log, cinder/volume.log
- Network
 - neutron/server.log

Debug journey - logs in Libvirt/Qemu

- Set the filters, log priorities for libvirtd
 - Add to file: /etc/libvirt/libvirtd.conf
 - log_level = 1
 - log_filters = "1:qemu 1:libvirt 1:json 1:monitor 3:event 1:util 1:security"
 - log_outputs = "1:[file:/var/log/libvirt/libvirtd.log](file:///var/log/libvirt/libvirtd.log)"
 - Restart libvirtd service
- Check the /var/log/libvirt/qemu/\$vm.log
 - libvirt-generated QEMU command-line arguments
 - QEMU error messages; libvirt stderr

Debug journey - auxiliary tools

➤ Virsh commands

- virsh start, virsh migrate, virsh list --all, virsh dumpxml
- service libvirtd status

➤ SR-IOV related commands

- ethtool, lsmod, modprobe, lspci -nnn, ip l set # up

➤ Numa/ hugepage related commands

- numactl --hardware, grep Huge /proc/meminfo
cat /sys/devices/system/node/node\$i/hugepages/...

➤ ABRT/ GDB

augtool, abrt-cli ls

What's next

- Get the alarm and proposed solution after the failures in OpenStack



Stack Tach ?

openstack.®

Call for actions

- Collect and analyze possible scenarios of VM failure and related OpenStack/ libvirt logs
- Propose solutions implementation in StackTach
- Raise more good ideas.....

Reference

- <https://access.redhat.com/documentation/en/red-hat-openstack-platform/>
- <https://docs.openstack.org/>
- <https://my.oschina.net/u/1179767/blog/752233>
- https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Deployment_Guide/ch-abrt.html



Thank you!

Chenli Hu
chhu@redhat.com