

# Much Ado About Blocking: Wait/Wake in the Linux Kernel

**LinuxCon China – June 2017. Beijing, China.**

**Davidlohr Bueso <dave@stgolabs.net>  
SUSE Labs.**



# Agenda

## 1. Fundamentals of Blocking

## 2. Flavors of blocking

- Wait-queues
- Simple wait-queues
- Lockless wake-queues
- rcuwait

# Fundamentals of Blocking

# Blocking 101

- As opposed to busy waiting, sleeping is required in scenarios where the wait time can be too long.
  - Other tasks ought to therefore run instead of wasting cycles.
  - Overhead of context switch vs wasting cycles.

# Blocking 101

- As opposed to busy waiting, sleeping is required in scenarios where the wait time can be too long.
  - Other tasks ought to therefore run instead of wasting cycles.
  - Overhead of context switch vs wasting cycles.
- `TASK_{UNINTERRUPTIBLE, INTERRUPTIBLE, KILLABLE}`
- There are three elements to consider when waiting on an event:
  - The wakee (`wake_up`, `wake_up_process`, etc).
  - The waker (`schedule`, `schedule_timeout`, `io`, etc).
  - The condition/event

# sleep\_on()

CPU0

```
while (!cond)
    sleep_on()
```

CPU1

```
cond = true
wake_up()
```

# sleep\_on()

CPU0

```
while (!cond)
```

```
    sleep_on()
```

CPU1

```
cond = true
```

```
wake_up()
```

- Inherently racy on SMP.
  - Missed wakeups
- Synchronization must be provided.
  - Locking
  - Memory barriers

# SMP-safe Blocking

```
for (;;) {  
    set_current_state();  
    if (cond)  
        break;  
    schedule();  
}  
__set_current_state(TASK_RUNNING);
```



# SMP-safe Blocking

```
for (;;) {  
    set_current_state();  
    if (cond)  
        break;  
    schedule();  
}  
__set_current_state(TASK_RUNNING);
```

Pairs with barrier in  
wake\_up()

Wait-queues

# Wait-queues

- Provide different wrappers, such that users do not suffer from races aforementioned.

```
for (;;) {  
    long __int = prepare_to_wait_event(&wq, &__wait,  
    state);  
  
    if (condition)  
        break;  
  
    schedule();  
}  
finish_wait(&wq, &__wait);
```

# Wait-queues

- Provide different wrappers, such that users do not suffer from races aforementioned.

```
for (;;) {  
    long __int = prepare_to_wait_event(&wq, &__wait,  
    state);  
    if (condition)  
        break;  
    schedule();  
}  
finish_wait(&wq, &__wait);
```

# Wait-queues

- Provide different wrappers, such that users do not suffer from races aforementioned.

```
for (;;) {  
    long __int = prepare_to_wait_event(&wq, &__wait,  
    state);  
  
    if (condition)  
        break;  
  
    schedule();  
}  
finish_wait(&wq, &__wait);
```

# Wait-queues

- Provide different wrappers, such that users do not suffer from races aforementioned.

```
for (;;) {  
    long __int = prepare_to_wait_event(&wq, &__wait,  
    state);  
    if (condition)  
        break;  
    schedule();  
}  
finish_wait(&wq, &__wait);
```

removal keeps the process  
from seeing multiple wakeups

# When Wait-queues Were Simple

- Basic linked list of waiting threads.
- A `wake_up()` call on a wait queue would walk the list, putting each thread into the runnable state.



# Wait-queues

- Exclusive Wait
  - Only the first N tasks are awoken.
- Callback mechanism were added for asynchronous I/O facility could step in instead.



# Wait-queues

- Lockless `waitqueue_active()` checks

CPU0 - waker

```
cond = true;
smp_mb();
if (waitqueue_active(wq))
    wake_up(wq);
```

CPU1 - waiter

```
for (;;) {
    prepare_to_wait(&wq, &wait, state);
    // smp_mb() from set_current_state()
    if (cond)
        break;
    schedule();
}
finish_wait(&wq, &wait);
```

# Wait-queues as Building Blocks

- Completions
  - Allows processes to wait until another task reaches some point or state.
  - Similar to `pthread_barrier()`.
  - Documents the code very well.
- Bit waiting.

# Wait-queues and RT

- Waitqueues are a big problem for realtime as they use spinlocks, which in RT are sleepable.
  - Cannot convert to raw spinlock due to callback mechanism.
  - This means we cannot perform wakeups in IRQ context.

# Simple Wait-queues (swait)

# Simple wait-queues

- Similar to the traditional (bulky) waitqueues, yet guarantees bounded irq and lock hold times.
  - Taken from PREEMPT\_RT.
- To accomplish this, it must remove wq complexities:
  - Mixing INTERRUPTIBLE and UNINTERRUPTIBLE blocking in the same queue. Ends up doing  $O(n)$  lookups.
  - Custom callbacks (unknown code execution).
  - Specific task wakeups (maintaining list order is tricky).

# Simple wait-queues

- Similar to the traditional (bulky) waitqueues, yet guarantees bounded irq and lock hold times.
  - Taken from PREEMPT\_RT.
- To accomplish this, it must remove wq complexities:
  - Mixing INTERRUPTIBLE and UNINTERRUPTIBLE blocking in the same queue. Ends up doing  $O(n)$  lookups.
  - Custom callbacks (unknown code execution).
  - Specific task wakeups (maintaining list order is tricky).
- Results in a simpler wq, less memory footprint.

# Simple wait-queues

- `swait_wake_all()`: task context

```
raw_spin_lock_irq(&q->lock);  
list_splice_init(&q->task_list, &tmp);  
while (!list_empty(&tmp)) {  
    curr = list_first_entry(&tmp, typeof(*curr), task_list);  
    wake_up_state(curr->task, TASK_NORMAL);  
    list_del_init(&curr->task_list);  
  
    if (list_empty(&tmp))  
        break;  
  
    raw_spin_unlock_irq(&q->lock);  
    raw_spin_lock_irq(&q->lock);  
}  
raw_spin_unlock_irq(&q->lock);
```

# Simple wait-queues

- `swait_wake_all()`: task context

```
raw_spin_lock_irq(&q->lock);  
list_splice_init(&q->task_list, &tmp);  
while (!list_empty(&tmp)) {  
    curr = list_first_entry(&tmp, typeof(*curr))  
    wake_up_state(curr->task, TASK_NORMAL);  
    list_del_init(&curr->task_list);  
  
    if (list_empty(&tmp))  
        break;  
  
    raw_spin_unlock_irq(&q->lock);  
    raw_spin_lock_irq(&q->lock);  
}  
raw_spin_unlock_irq(&q->lock);
```

drop the lock (and  
IRQ disable)  
after every wakeup



# Lockless Wake Queues

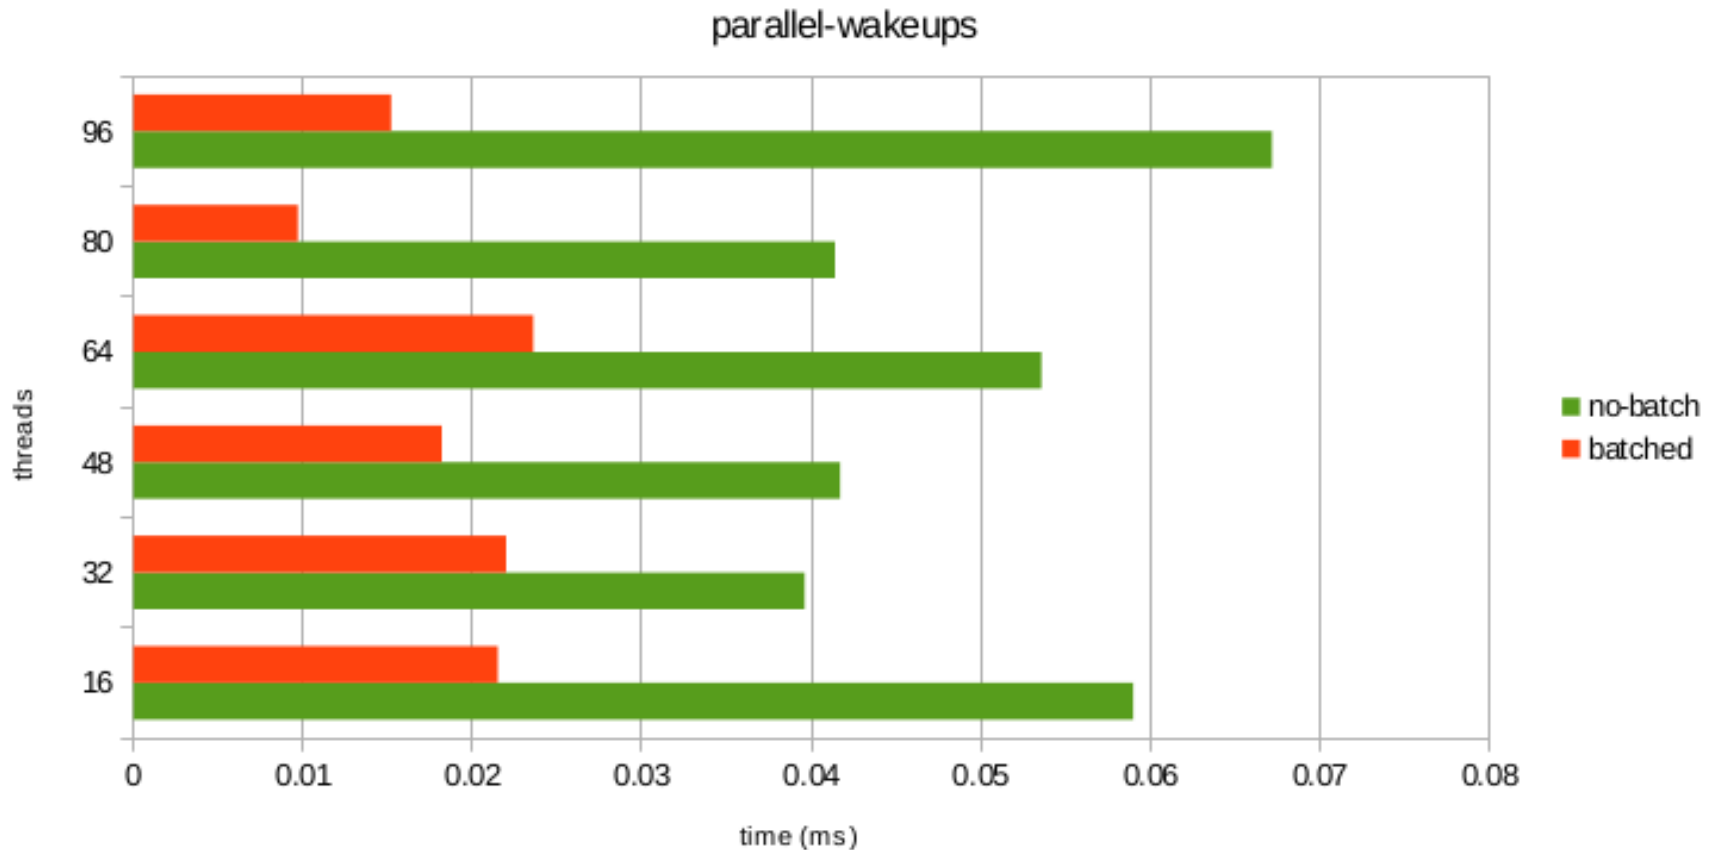
# Wake Queues

- Internally acknowledge that one or more tasks are to be awoken, then call `wake_up_process()` after releasing a lock.
  - `wake_q_add()/wake_up_q()`
- Hold reference to each task in the list across the wakeup thus guaranteeing that the memory is still valid by the time the actual wakeups are performed in `wake_up_q()`.

# Wake Queues

- Maintains caller wakeup order.
- Works particularly well for batch wakeups of tasks blocked on a particular event.
  - Futexes, locking, ipc.

# Wake Queues



26-core, 2 socket x86-64 (Haswell)

rcuwait

# rcuwait

- `task_struct` is not properly rcu protected unless dealing with an rcu aware list
  - `find_task_by_*()`.
- `delayed_put_task_struct()` (via `release_task()`) can drop the last reference to a task.
  - Bogus wakeups, etc.
- Provides a way of blocking and waking up a single task in an rcu-safe manner.



# rcuwait

- But what about `task_rcu_dereference()`?
  - Task freeing detection
  - `probe_kernel_read()`
  - May return a new task (false positives)



# rcuwait

- If we ensure a waiter is blocked on an event before calling `do_exit()/exit_notify()`, we can avoid all of `task_rcu_dereference` overhead and magic.
- Currently used in percpu-semaphores.



# References

# References

- Documentation/memory-barriers.txt
- Documentation/scheduler/completion.txt
- [Simple Wait Queues LWN](#)
- [Return of Simple Waitqueues \(LWN\)](#)
- Source Code:
  - kernel/sched/{wait,swait}.c
  - kernel/sched/exit.c (rcuwait)
  - include/linux/sched/wake\_q.h

Thank you.

