

Unikernelized Linux

Tiejun Chen <tiejunc@vmware.com>

VMware China R&D Advanced Technology Center



Warning 😊

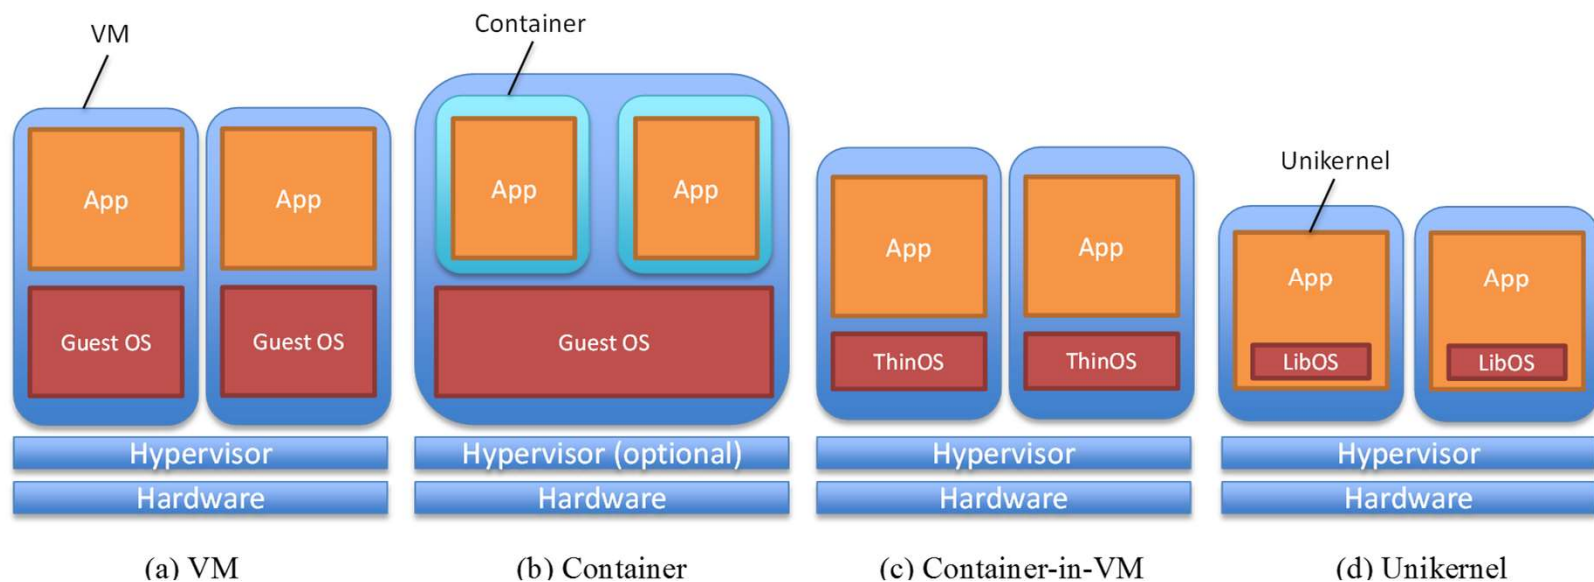
This is our own exploration of unikernels.

This is not a roadmap or commitment from VMware.



Background

- **Linux container technologies like Docker dominate dramatically**
 - An efficient but easy way to carry out applications to provide cloud services in the different cases.
- **A new technology called unikernels is beginning to attract our attention**
 - Unikernels are developing a variety of new approaches to deploy cloud services.



Major Existing Unikernels

- **Unikernels projects**

- MirageOS, ClickOS, Clive, HaLVM, LING, Rump Kernels, OSv, Unik, Solo5 Unikernel, Drawbridge

- **Unikernels solutions**

- Docker
 - Hyperkit/VPNkit
 - Moby/Linuxkit
- Mikangelo
 - Improving Responsiveness and Agility of HPC Cloud Infrastructure
- NFV
 - Unikernels based NFV architecture



Unikernel Definition and Types

- **Definition**

- Unikernels are specialised, single address space machine images constructed by using library operating systems. --Wiki

- **Types**

- General purpose unikernels
 - A library that derives from a generally designed OS kernel
 - Works for apps that follow some mature speculations (e.g. POSIX, or glibc)
 - Example: Rumprun, OSv, ClickOS and Drawbridge
- Language specific unikernels
 - A library of a programming language that includes all OS functionalities
 - Works for apps written in specific languages only
 - Example: MirageOS (OCaml), Clive (Golang), HalVM (Haskell), IncludeOS (C++), Ling (Erlang) and Runtime.js (Javascript)



Unikernel Essentials

- **The biggest characteristics**
 - Single address space: Zero-copy and huge page
 - Single running mode: Perform the efficient function call
 - One process with multiple threads: No heavy context switch and TLB flush
- **Compared to a traditional OS, unikernels provide many benefits**
 - Improved security
 - Unikernels reduce the amount of code deployed, which reduces the attack surface.
 - Small footprints
 - Unikernels images are often orders of magnitude smaller than traditional OS deployments.
 - Highly optimized
 - Unikernels enables whole-system optimization across device drivers and application logic. Especially it's mostly paravirtualized under virtualization environment.
 - Fast Boot
 - Unikernels can boot extremely quickly, with boot times measured in milliseconds.

Unikernel Experiments: Public Claims

- **OSv**

- For unmodified network-intensive applications, we demonstrate up to 25% increase in throughput and 47% decrease in latency. By using non-POSIX network APIs, we can further improve performance and demonstrate a 290% increase in Memcached throughput.
- <http://www.cs.utah.edu/~peterm/prelim-osv-performance.pdf>

- **IncludeOS**

- As a test case a bootable disk image consisting of a simple DNS server with OS included is shown to require only 158 kb of disk space and to require 5-20% less CPU-time, depending on hardware, compared to the same binary running on Linux.

- **ClickOS**

- ClickOS virtual machines are small (5MB), boot quickly (about 30 milliseconds), add little delay (45 microseconds) and over one hundred of them can be concurrently run while saturating a 10Gb pipe on a commodity server.



Unikernel Challenges

- **Challenges - why existing unikernels have yet to gain large popularity**
 - Lack of compelling use cases
 - Compatibility with existing applications
 - Lack of production support (e.g. monitoring, debugging, logging)



Use Cases for Unikernels 1/3

- **Serverless**

Most public cloud vendors are embracing this promising model with container.

- Pros

- Quick OS boot & improved security & smaller size and footprint
- Mature VM management
- Potentially multiple languages support

- Cons

- Unikernels is a little heavy to carry out just one function.
- Debug issue can worsen serverless development.
- Time of creating VM has a significant impact on function invocation.

- Conclusion

- In terms of different QOS unikernels are beneficial and useful complement to serverless mode. Furthermore, what if we can unikernelize linux, and further optimize it accordingly.



Use Cases for Unikernels 2/3

- **IoT**

IoT is a big markets as well.

- Pros

- The feature of smaller size & footprint are good for those resource-strained IoT platforms.
- Such a lightweight VM instance can address security issue.

- Cons

- Oftentimes unikernels need virtualization technology.
- Unikernels are not designed to address those IoT characters like power consumption.
- Unikernels don't support versatile architectures.

- Conclusion

- Unikernels can value IoT when virtualization probably thrives at the edge. More importantly, IoT closely ties with the embedded system where Linux always plays a very import role, so it's worth fitting unikernlized Linux into IoT.



Use Cases for Unikernels 3/3

- **IO intensive applications**

IO Performance always captures people's attention.

- Pros

- Oftentimes unikernels have the simple IO flow framework

- Cons

- Only a subset of I/O intensive apps are good for unikernels: the latency-sensitive apps. The other subset of I/O intensive apps like the bandwidth-intensive apps need more considerations and explorations.

- Conclusion

- Unikernels can contribute IO case at large. NFV is really a potential chance to make unikernels succeed with any targeted acceleration to Linux.



Exploration Conclusions

- **Summary**
 - Unikernels still yield comparable performance.
 - The different unikernels have different focuses.
 - User has to put more or less effort to develop an application based on unikernels.
 - Nothing more specifically is done to embrace unikernels from hypervisor's view.
 - Linux could be a good candidate of unikernels
 - Linux itself could help eliminate those challenges of unikernels
 - All optimizations and acceleration aimed to Linux can benefit unikernels
 - Unikernelized Linux can catch more eye by means of Linux community 😊



What Could We Do?

Our target is to explore what is the best platform for running unikernels case


We will achieve this by

- **Research existing unikernels**
 - Integrate and support those major existing unikernels well
- **Build new unikernel**
 - Convert Linux kernel
- **Explore optimizations**
 - Integrate virtIO model into ESXi as an example
 - Provide monitoring, logging and remote debugging
 - Supporting a short lived unikernels instance
 - Resources are consumed by live unikernels



What Are The Key Challenges?

- **Convert Linux to unikernels**
 - The fundamental philosophy of Linux is aiming to multiple processes and two modes.
 - Most components are coupled tightly.
 - How to further improve performance
- **Reduce time of creating VM**
 - Snapshot
 - VM Fork
- **A good paravirtualized API for common unikernels**
 - Some pv ops might already be a good start
- **New scheduler**
- **Manage the lifecycle and identities of the provisioned unikernels**



How Could We Possibly Achieve This? Hypervisor basics

- **Support major existing unikernels**
 - Integrate virtIO framework into ESXi
 - Port vmxnet3 and pvscsi into them
- **Define a standard API which can paravirtualize unikernels**
 - Based on common hypercall
 - Configure/control guest OS
 - Setup Inter-VM Communication
 - Allocate/destroy memory directly
- **Add a new scheduler**
 - Address short lived VM
 - Schedule a group of unikernels instances



How Could We Possibly Achieve This? Linux basics 1/2

- **Convert Linux**
 - Single running mode
 - Ring 0
 - `__USER32_CS | __USER_DS | __USER_CS`
 - Check with `'cmpq $__USER_CS, CS(%rsp)'`
 - Stack
 - Switch stacks manually
 - Interrupt Stack Table (IST)
 - `set_intr_gate_ist(X86_TRAP_PF, &page_fault, PF_STACK)`
 - Interrupt and exception
 - Single address space
 - Single process
 - No `fork()/exec()`



How Could We Possibly Achieve This? Linux basics 2/2

- **Convert Linux**
 - Optimization
 - Smaller size and footprint
 - Zero-copy
 - {get,put}_user
 - copy_{from,to}_user
 - Other unnecessary copy and check
 - Scheduler
 - scheduling classes & policies
 - fair vs rt vs deadline
 - New?
 - Lightweight TCP/IP Stack
 - LWIP
 - Fastsocket
 - Seastar
 - A variety of Linux variants
 - Multiple Unikernelized Linux profiles



How Could We Possibly Achieve This? Compatibility

- **Support existing applications**
 - Different code circumstances
 - Source code
 - New standard library
 - glibc
 - Function Call
 - Binary
 - `–shared –pic`
 - `LD_PRELOAD`
 - Others
 - BT
 - Multiple processes
 - One fork = one unikernelized Linux instance
 - IPC = Inter-VM Communication
 - PCID – Process-context identifiers
 - Limited bits
 - Linux's own debug/monitor/log tools and utilities



How Could We Possibly Achieve This?

Debugging, monitoring and logging

- **Debug unikernels**
 - Log info
 - virtual serial port
 - Dynamic buffer memory allocation
 - Linux's own utilizes
 - ssh/gdb/ftrace/perf/kprobe/kdump/...
 - PCID & the balloon driver
- **Monitor unikernels**
 - A mini-httpd as a stub connecting those Linux utilities
 - Inspired by OSv
- **Log unikernels**
 - rsyslog
 - vRealize Log Insight



How Could We Possibly Achieve This? Enhancements

- **Offer faster boot**
 - Explore ESXi to further reduce the time of creating VM
 - Skip BIOS with a small integrated bootloader
 - Replace ACPI with DTB
 - Adopt 1:1 Bus/device initialization
 - No any redundant bus scanning and device probing
- **Utilize hardware virtualization**
 - VT-X Instructions
 - VMFUNC
 - Pre-construct EPT table to get a faster and secure way to communicate between unikernels
 - VT-X Features
 - VPID (Virtual processor ID)
 - The tagged TLB to reduce cost of performance
 - Preempt Timer
 - A feature which count down in unikernels without too much external timer injected by hypervisor

How Could We Possibly Achieve This? Others

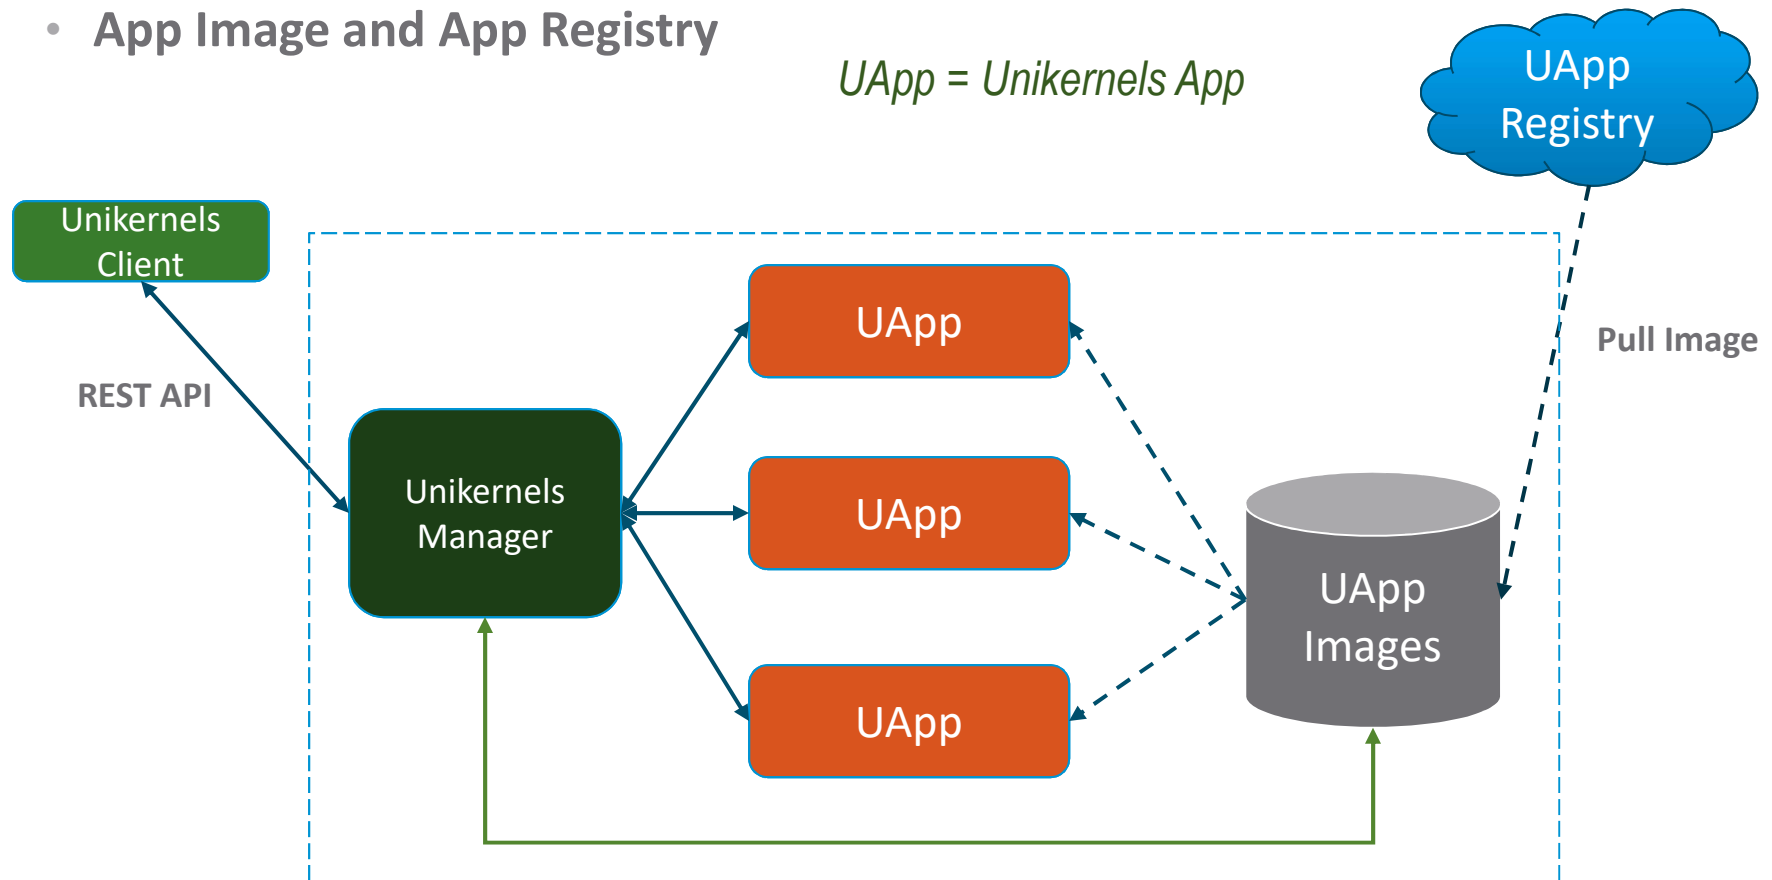
- **Construct an efficient toolchain**
 - Build and deploy unikernels like Docker
 - Customized components management
 - Configuration
 - Kernel image
 - User App
 - Dependencies
- **Support orchestration**
 - Docker Swarm Mode, Kubernetes, Mesos and Cloud Foundry
 - **Unik**
- **Integrate Source Code Analyzer tool**
 - This can help us enhance security from code level



How Could We Possibly Achieve This? Management

- Unikernels Manager
- App Image and App Registry

UApp = Unikernels App



References

- <http://unikernel.org/projects/>
- [https://wiki.xen.org/images/3/34/XenProject Unikernel Whitepaper 2015 FINAL.pdf](https://wiki.xen.org/images/3/34/XenProject_Unikernel_Whitepaper_2015_FINAL.pdf)
- <https://www.linux.com/news/7-unikernel-projects-take-docker-2015>
- <https://www.usenix.org/node/184012>
- <https://www.deepdyve.com/lp/institute-of-electrical-and-electronics-engineers/includeos-a-minimal-resource-efficient-unikernel-for-cloud-services-J43NrZQ7fn>
- <https://www.usenix.org/system/files/conference/nsdi14/nsdi14-paper-martins.pdf>



Thank You!

tiejunc@vmware.com

