

GPU Acceleration for Container on Intel Processor Graphics

Zhenyu Wang
zhenyuw@linux.intel.com

Agenda

GPU Containers

GPU Namespace


GPU Control Group

Integration with Container Runtime

VM Containers

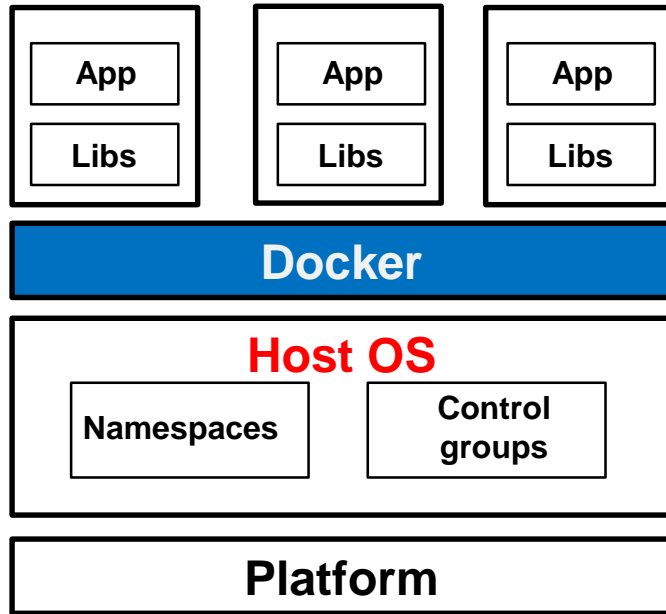
Status

GPU Containers




What Makes a Container?

Bare Metal Containers



VM Containers (e.g. Clear Container)



Path to GPU Containers

Bare Metal Containers

GPU namespaces

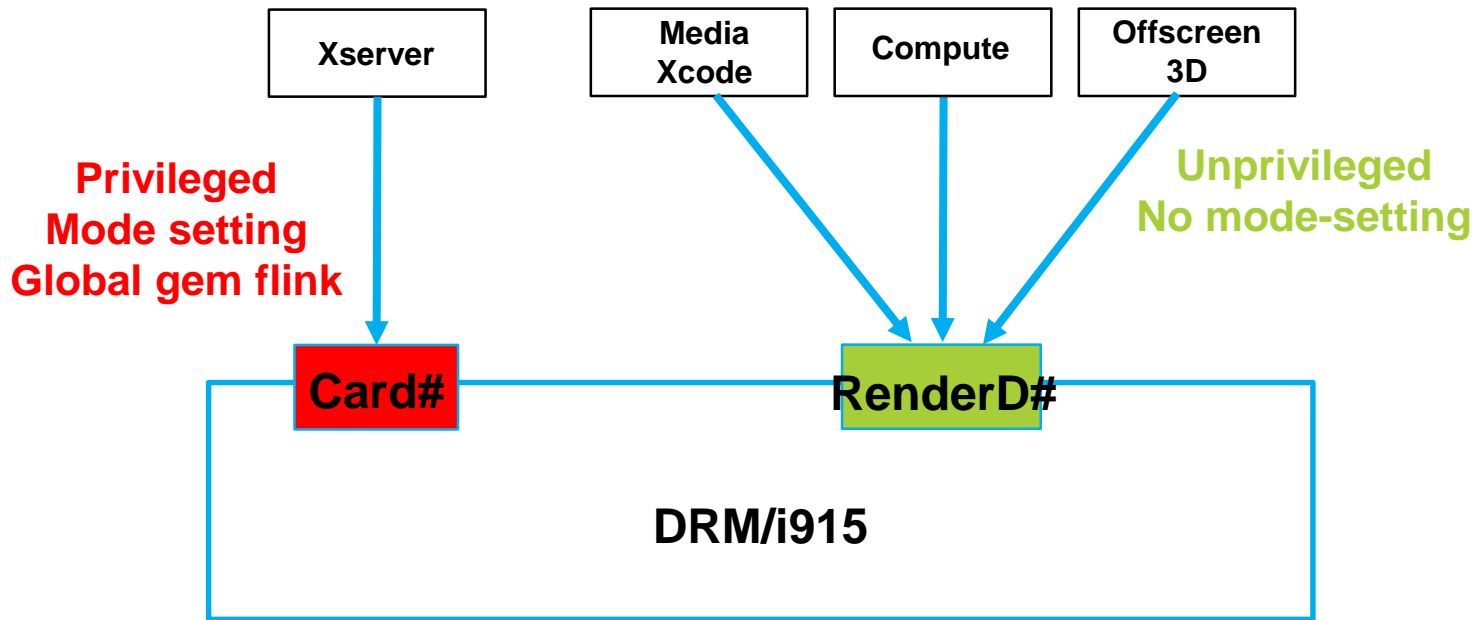
GPU control groups

VM Containers

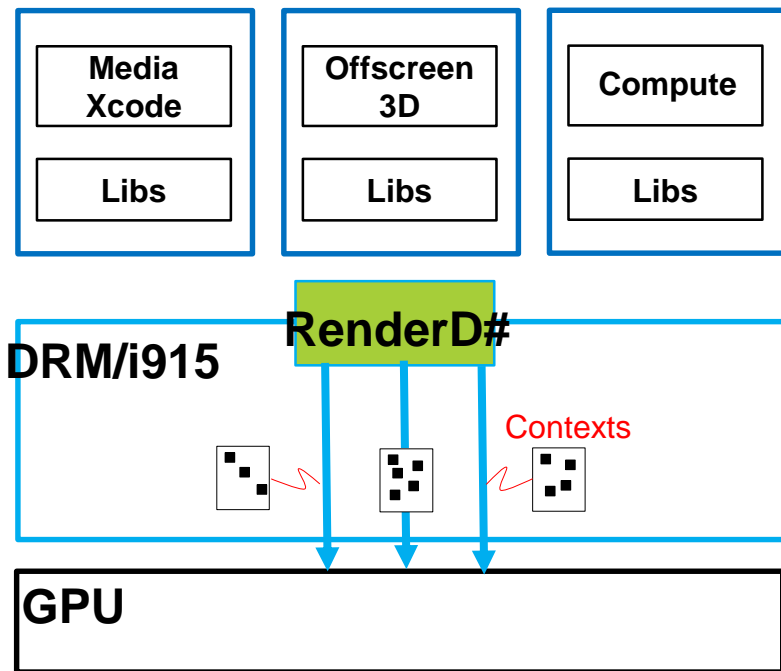
GPU virtualization technology

Integrate with container runtime

GPU Namespace: DRM Render Node



GPU Namespace: DRM Render Node



Sharing happen with dmabuf only (no global gem object on card0)

No need of creating another namespace

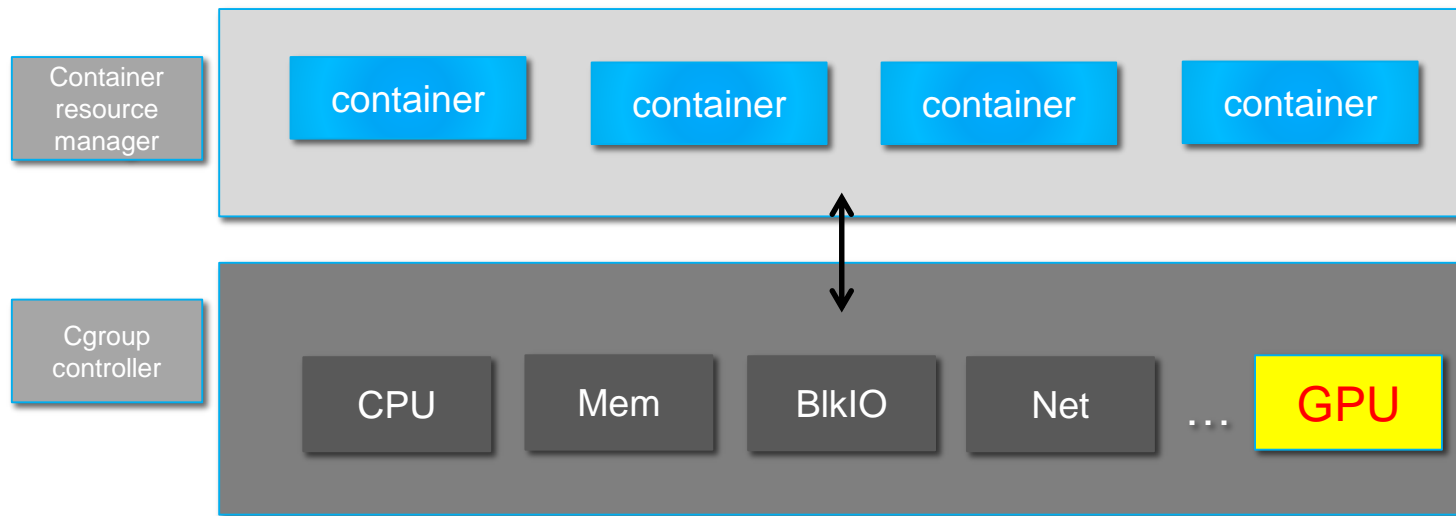
Expose DRM render nodes to provide isolated views through device cgroup

- ✓ e.g “`docker run --device=/dev/dri/renderD128 --it debian`”

Per-process hardware context in GPU

- ✓ Hardware managed render context switch

GPU CGroup



Container-granule GPU resource control

GPU CGroup

New subsystem: `gpu_cgrp_subsys`

Control knobs

- ✓ `gpu.shares` (% share of GPU cycles, work in progress)
- ✓ `gpu.priority` (workload priority in the system)
- ✓ `gpu.memory` (maximum GPU memory size)

DRM/i915

- ✓ Favor 'shares/priority' in workload scheduler
- ✓ Enforce memory limitation (optionally for UMA graphics)

Container Runtime Integration

RunC: an universal container runtime

- ✓ Understand new gpu cgroup
- ✓ Add new Linux resources config.json options for gpu
- ✓ Runtime-tools helper to generate config stubs

New Docker GPU control parameters

e.g `docker -gpu-mem=xxx -gpu-priority=xxx
-gpu-share=xxx ...`

```
{  
  "linux": {  
    "devices": [  
      {  
        "path": "/dev/dri/renderD128",  
        "type": "c",  
        "major": 226,  
        "minor": 128,  
        "fileMode": 432,  
        "uid": 0,  
        "gid": 0  
      }  
    ],  
    "resources": {  
      ...  
      "gpu": {  
        "memory": <max_mem_in_bytes>,  
        "prio": <workload_priority>  
      }  
    }  
  }  
}
```

Image Portability

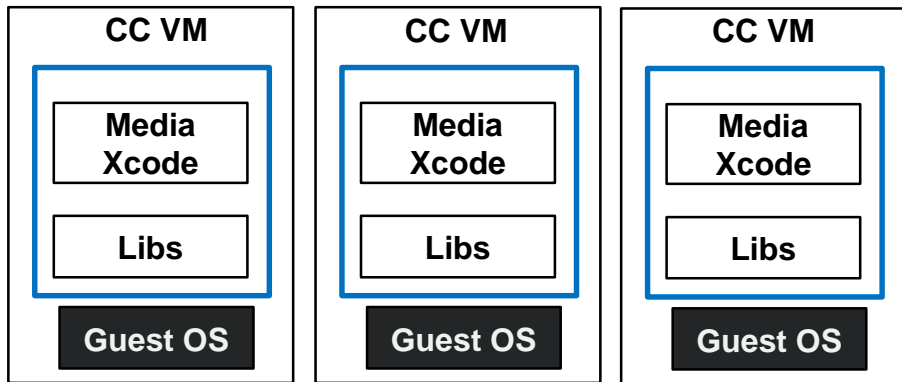
Main challenge - library compatibility

- ✓ User level libraries MUST match underlying kernel/HW

Introduce Docker helper for image inspection

- ✓ Compare image libraries with host environment
- ✓ Reject to launch container upon any mismatch

VM GPU Containers



DRM/i915

Intel GVT-g

vGPU

vGPU

vGPU

GPU

Clear Container (CC)

Intel graphics virtualization technology
(Intel GVT-g)

Assign vGPU to VM container

In upstream Linux since v4.10

<https://01.org/igvt-g>

VM GPU Containers

Nested containers in VM also have GPU access

Need Improvement:

Scalability

- ✓ Only support 8vGPUs today due to resource partitioning
- ✓ Need some enlightened way to further scale in the short term

Boot time

- ✓ Full guest graphics driver load takes >1.5s
- ✓ Fast optimization to <0.5s in progress

Status

- ✓ GPU cgroup PoC

<https://github.com/zhenyw/linux>

<https://github.com/zhenyw/runc>

<https://github.com/zhenyw/moby>

<https://github.com/zhenyw/cli>

- ✓ GPU virtualization for Clear container

<https://clearlinux.org/features/intel%C2%AE-clear-containers>

<https://github.com/01org/gvt-linux/wiki/Clear-Container-with-GVTg-Setup-Guide>

Q & A